Binary Search Trees

Dr. Anirban Ghosh

School of Computing
University of North Florida

Maps

Definition

A **record** is a key-value pair: (k, v)

A **map** is an abstract data type for maintaining a set of records

- No two records cannot have the same key
- However, two records can have same values though
- Association of keys to values define a **mapping**: $f(\text{key}) = \texttt{value}$

Examples of maps

- UNF maintains a map of `(N#, student information)` records
- A social media company maintains a map of `(email address, user account information)` records
- An assembler maintains a symbol table (a map) of `(opcode, hex)` records
- A text-editor maintains a map of `(color, RGB representation)` records

How to implement a map?

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k

Approach 1: maintain a sorted list of records

- **Insertion.** will take $O(n)$ time for figuring out the correct spot for the record; then $O(n)$ time for shifting items to the right to accommodate the new record; total time taken is $O(n)$
- **Retrieval.** will take $O(\log n)$ time using binary search
- **Deletion.** will take $O(\log n)$ time to locate it using a binary search; then then $O(n)$ time for left shifting items to kill the empty spot; total time taken $O(\log n) + O(n) = O(n)$

How to implement a map?

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k

Approach 2: maintain an unsorted list of records

- **Insertion.** will take $O(1)$ time (add at the end)
- **Retrieval.** will take $O(n)$ time using a linear search; may need to search the whole list in the worst-case
- **Deletion.** will take $O(n)$ time to locate it using a linear search; then $O(n)$ time for left shifting items to kill the empty spot; total time taken $O(\log n) + O(n) = O(n)$

Our aim

Common map operations

- **Insert** a record (k, v)
- **Retrieve** a record having key k
- **Delete** a record having key k

☞ To accomplish the above three tasks in $O(\log n)$ time each

☞ **Balanced** binary search trees is the solution; stay tuned ...

The map ADT

```java
public interface MapADT<K,V> {
   boolean put(K key, V value); // adds a new record with key 'key' and value 'value'
   V remove(K key); // removes the record having key 'key'
   V get(K key); // return the value part of the record whose key is 'key'
   V updateValue(K key, V newValue); // updates the value part of the record whose key is 'key' with a new value
   int size(); // returns the number of records stored in the map
   void clear(); //Removes all records from the map
}
```

What is a Binary Search Tree?

- It is a **binary tree** where every node contains a **<key, value>** pair (a **record**); keys must be **comparable** but the values don't need to be
- Moreover, for every node p in the tree, the following 2 properties hold
  1. Keys stored in the left subtree of p are $<$ the key stored at p
  2. Keys stored in the right subtree of p are $>$ the key stored at p

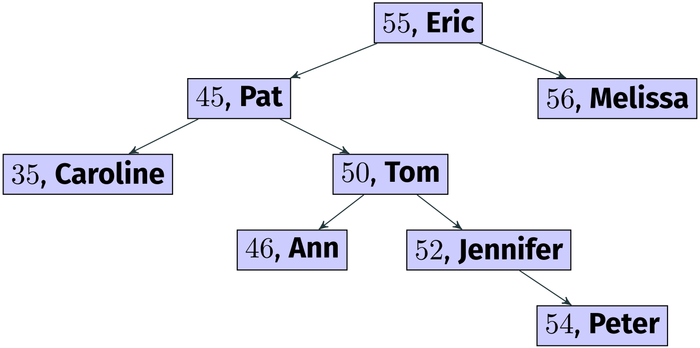

What is a Binary Search Tree?

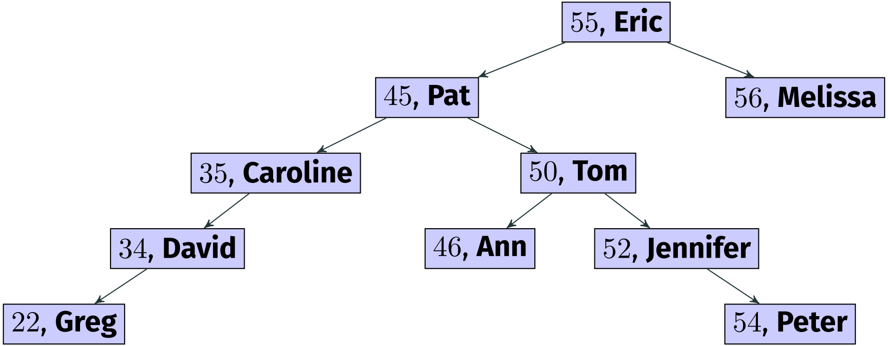

Use

BSTs can be used to implement **maps** and are commonly used for fast searching (typically need far less comparisons than lists)

An important property of BSTs

An inorder traversal of a BST always gives the sorted sequence based on the keys

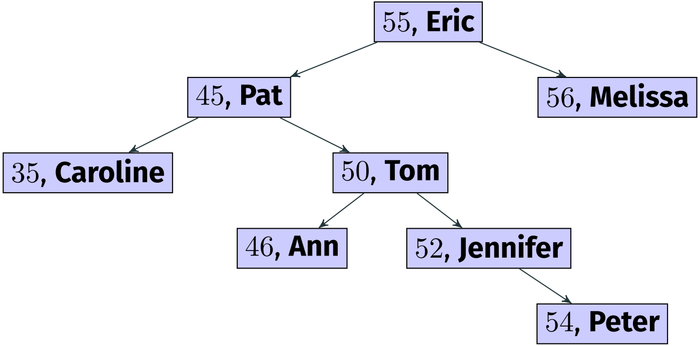

Inorder traversal

35, **Caroline**; 45, **Pat**; 46, **Ann**; 50, **Tom**; 52, **Jennifer**; 54, **Peter**; 55, **Eric**; 56, **Melissa**

Let's say you need to look for the record that has the key k; how will you do this?

Algorithm

- Start at the root
- If the root's key is k, then search is successful
- If $k <$ root's key, search recursively (or iteratively) in the left subtree of the root
- Otherwise, search recursively (or iteratively) in the right subtree of the root
- If we have reached a `null` link, no record exists in the tree with key k

To save space in the figures, we will write only the record keys inside the nodes and avoid the corresponding values

Example

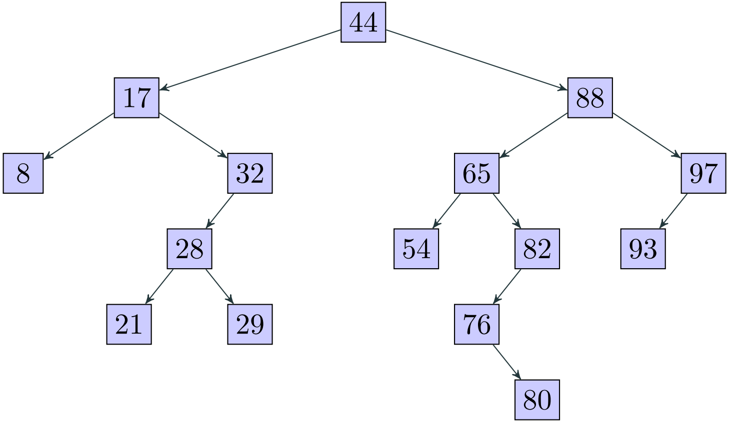

Example: search for 28

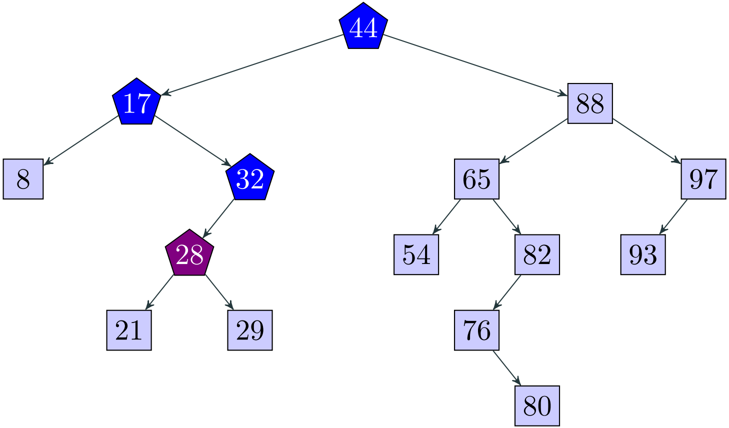

Found!

Example: search for 68

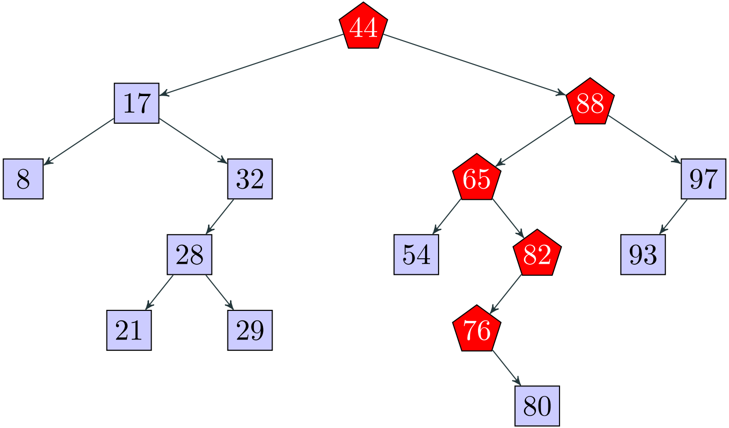

Not found; a null link is reached (the left link of 76)

$O(h)$, where h is the height

How to insert a record into a BST having key k?

Algorithm

- Start at the root
- If the root is `null`, replace empty tree with a new tree with the new record as the root, and signal **SUCCESS**
- If k equals root's key, signal **FAILURE** since a record with key k already exists
- If $k <$ root's key, insert recursively (or, iteratively) in the left subtree of the root
- Otherwise, insert recursively (or, iteratively) in the right subtree of the root

Example: insert 68

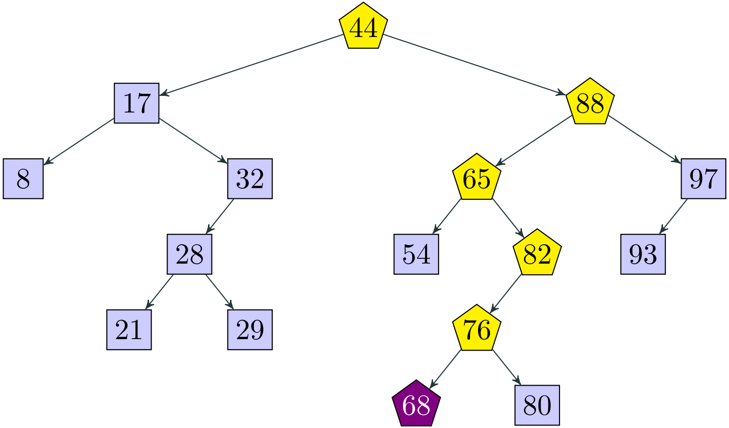

Time complexity of insertion

$O(h)$, where h is the height

How to delete the record from a BST having key k?

Idea

- Traverse through the tree to find the record having key k
- If the record cannot be found, no action is needed
- Now assume that we have found the record that has key k at node r
  1. r is a leaf node (no child)
  2. r has exactly one child (either left or right)
  3. r has two children (both left and right)

Case 1: r is a leaf node, delete 29

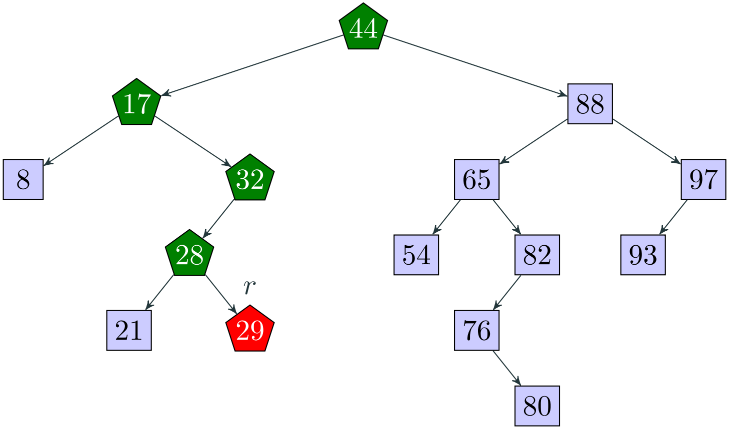

Easy! just delete it right away

Case 1: r is a leaf node, delete 29

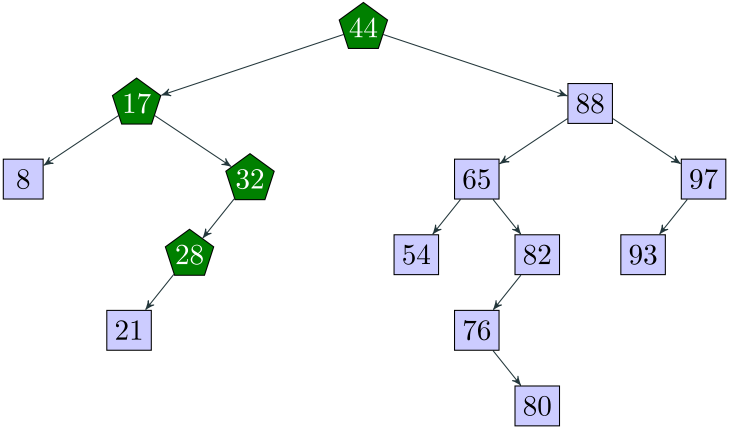

Case 2: r has one child, delete 97

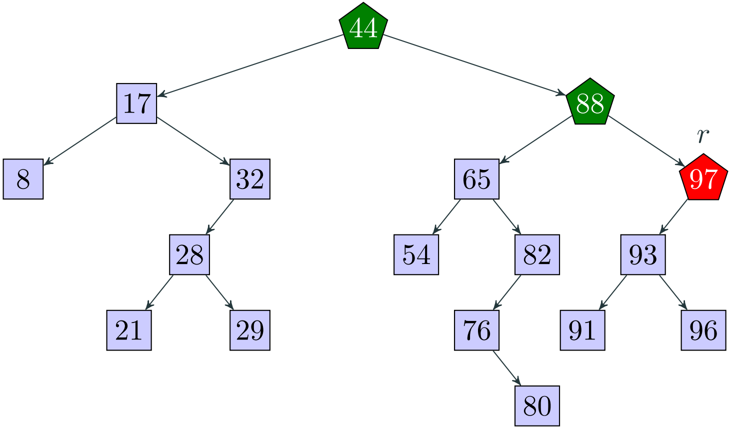

Make the parent of r point to the only child of r instead of r

Case 2: r has one child, delete 97

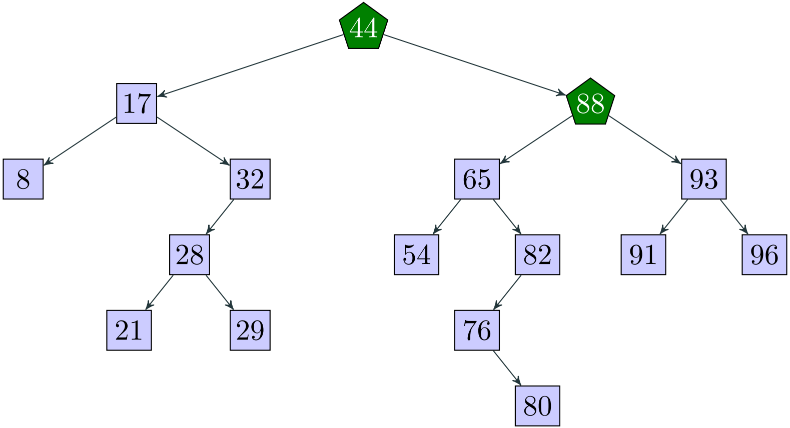

Case 3: r has two children, delete 88

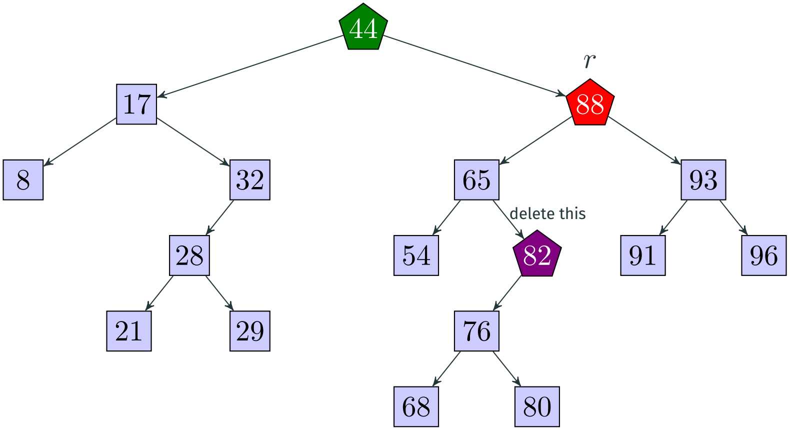

Replace the record at r with the record at the inorder predecessor p of n; then delete p using Case 1 or 2, depending on the number of children of p. Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Case 3: r has two children, delete 88

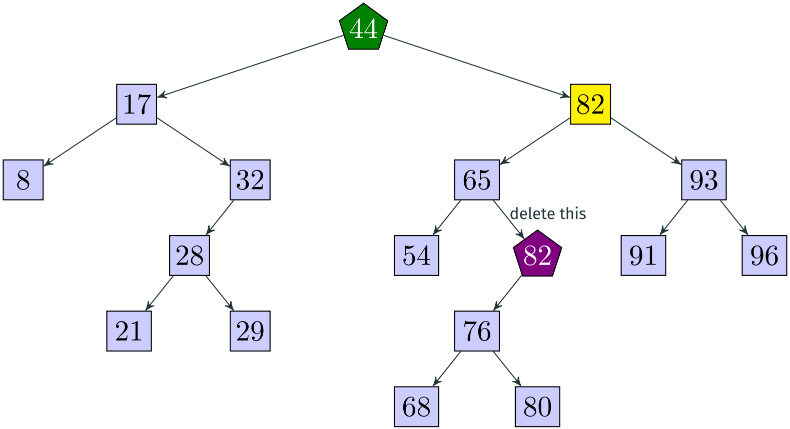

Replace the record at r with the record at the inorder predecessor p of n; then delete p using Case 1 or 2, depending on the number of children of p. Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Case 3: r has two children, delete 88

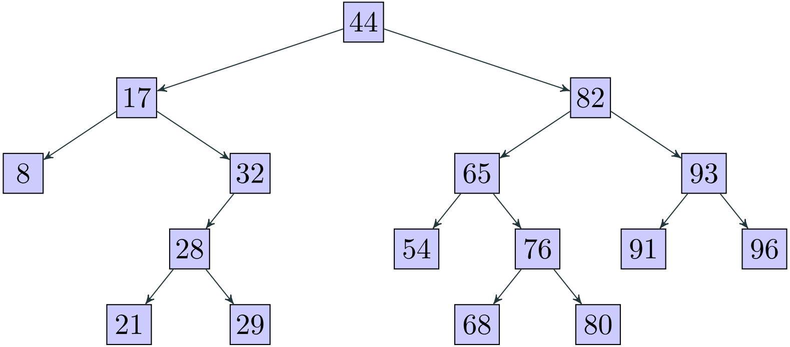

Replace the record at r with the record at the inorder predecessor p of n; then delete p using Case 1 or 2, depending on the number of children of p. Note that the inorder predecessor is either a leaf node or has a left child only (no right child).

Time complexity of deletion

$O(h)$, where h is the height

Code

See the class `TreeMapBST`

BSTs are not unique

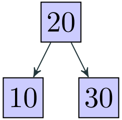

Insertion sequence: $20, 10, 30$

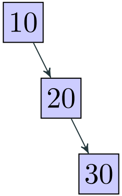

Insertion sequence: $10, 20, 30$

Their structures really depend on the insertion sequence of records

Problem

Given a text, find out the unique words in it along with their counts. We also need to output the distinct words with their count.

A solution

- Use a BST T where the key-type is `String` and the value-type is `Integer`
- For every word w in the text, first check if a node exists in T, where the stored key is w
  - If such a node does not exist, insert a new node in T with key w and value 1
  - If such a node exists in T, increment the stored value (essentially a counter) by 1

See the class `WordCounter`

Worst case scenario: skewed binary trees

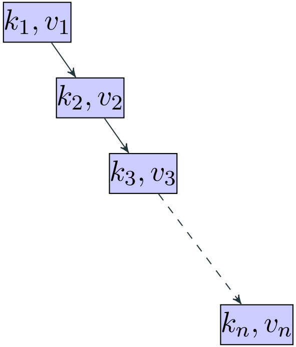

In this case, $k_1 < k_2 < k_3 < \ldots < k_n$

So, in the worst-case, $h = n - 1 = O(n)$

Therefore, searching, insertion, deletion take $O(h) = O(n)$ time each!

This is as bad as using singly linked-lists!

Do something so that h remains bounded
We aim for $h = O(\log n)$

A solution. use **Red-Black** trees

Red-Black trees are never skewed or close to being skewed like the plain binary search trees we just talked about 👍

Rotation on BSTs

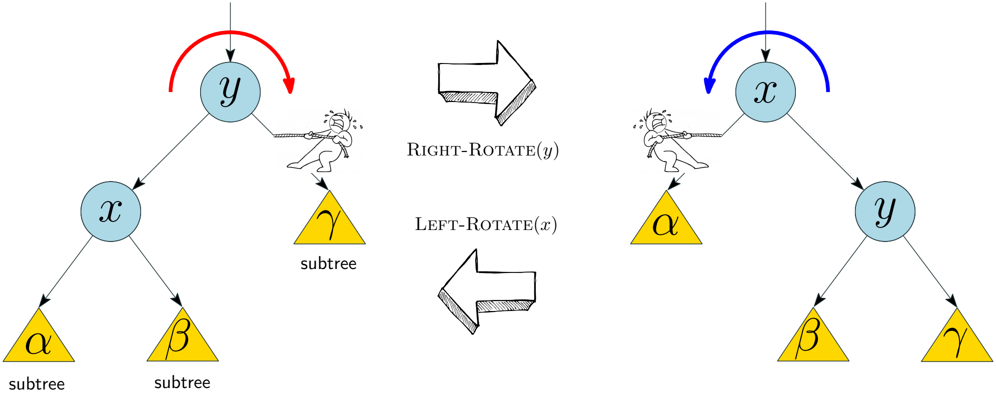

Rotations help in **reducing** height of BSTs; this means faster operations on BSTs

A single rotation can be done in $O(1)$ time

Can be applied to any type of **self-balancing** BST (red-black, AVL, etc.)

An example of right rotation

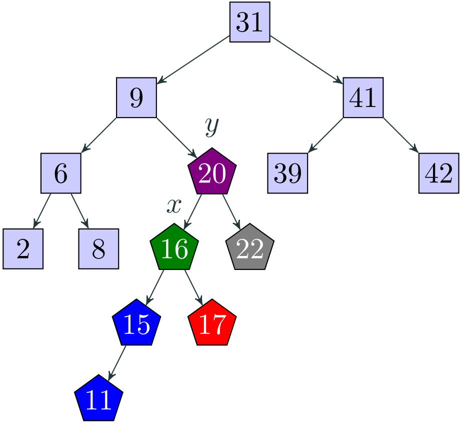

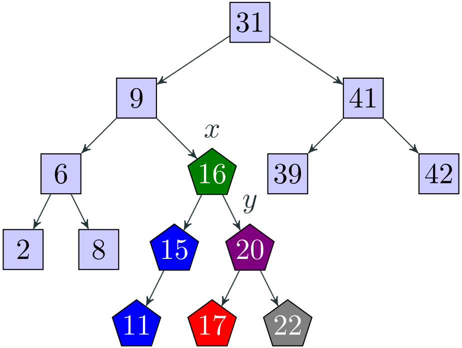

Height: 5

Height: 4

Note that the inorder traversal sequence remains the same!

Definition

A **red-black** tree is a self-balancing BST that maintains the following invariants:

1. Every node is either **red** or **black**
2. The root is **black**
3. If a node is **red**, then both its children are **black** (The `null` references of the leaves are **black**)
4. The number of **black** nodes in any path from the root to a leaf is the same

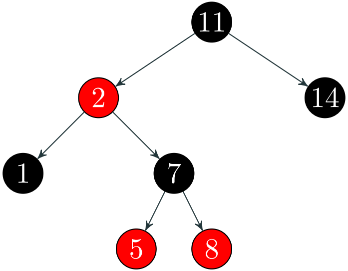

It can be shown mathematically (out of scope) that for red-black trees,

$$h = O(\log n)$$